

# Parallel-in-time integration with PFASST

## Swiss Numerical Analysis Day 2015

Daniel Ruprecht, Robert Speck, Rolf Krause

in collaboration with M. Emmett, M. Minion, M. Bolten and others

Institute of Computational Science  
Università della Svizzera italiana  
Lugano, Switzerland

17 April 2015

# Joint work with



Rolf Krause



Matthew Emmett  
Michael Minion



Matthias Bolten



Robert Speck

# A fundamental turn toward concurrency

- ▶ In about 10 years, **Exascale** systems will require 100-million way parallelism
- ▶ Multitude of challenges for numerical methods:
  - ① massive concurrency
  - ② resilience against faults
  - ③ energy consumption
  - ④ ...




Figure: Number of cores in the top system in [www.top500.org](http://www.top500.org).


*#8: Don't rethink your algorithms.*

(K. Yelick, "Ten ways to waste a parallel computer", 2009)

# 51 years of parallel-in-time integration<sup>1</sup>

## Parallelization in time can

- ① extend strong scaling limits of space parallelization
  - ② help with the "trap of weak scaling"
- Interpolation-based approach (Nievergelt 1964)
  - Parabolic or time multi-grid (Hackbusch 1984) and (Horton 1992)
  - Parallel Runge-Kutta methods (e.g. Butcher 1997) and extrapolation (Richardson 1910)
  - Parareal (Lions, Maday, Turinici 2001) and PITA (Farhat, Chandesris 2003)
  - PFASST (Emmett, Minion 2012)
  - MGRIT (Falgout et al. 2014)
  - DG time multi-grid (Gander, Neumueller 2014)



<sup>1</sup>For a recent overview see M. Gander's paper.

# Collocation



**Figure:** A time-step  $[T_n, T_{n+1}]$  with  $M = 9$  Gauss-Lobatto collocation nodes  $t_j$ .

**Integrate initial value problem  $u'(t) = f(u(t), t)$  from  $T_n$  to  $T_{n+1}$ :**

- Consider Picard formulation of IVP

$$u(T_{n+1}) = u(T_n) + \int_{T_n}^{T_{n+1}} f(\tau, u(\tau)) d\tau$$

- The integral is approximated by quadrature

$$\int_{T_n}^{T_{n+1}} f(\tau, u(\tau)) d\tau \approx \Delta t \sum_{j=1}^M q_{n,j} f(t_j, u_j), \quad \text{with stages } u_j \in \mathbb{R}^N$$

- Collocation problem can compactly be written as

$$\mathbf{U} = \mathbf{U}_0 + \Delta t \mathbf{QF}(\mathbf{U}), \quad \mathbf{U} = (u_1, \dots, u_M)^T \in \mathbb{R}^{NM \times NM}$$



# Spectral deferred corrections (SDC)



Dutt, Greengard, Rokhlin

BIT Numerical Mathematics 2000




Figure: Use composite rectangular rule  $\mathbf{Q}_\Delta$  to precondition high-order collocation rule  $\mathbf{Q}$

- Consider preconditioned Richardson iteration

$$\mathbf{Y}^{k+1} = \mathbf{Y}^k + (\mathbf{I} - \Delta t \mathbf{Q}_\Delta \mathbf{F})^{-1} [\mathbf{Y}_0 - (\mathbf{I} - \Delta t \mathbf{Q} \mathbf{F}) \mathbf{Y}^k], \quad \mathbf{Q}_\Delta \approx \mathbf{Q}, \quad \mathbf{Y} \in \mathbb{R}^{NM}$$

- Can compute this node-by-node which gives "classical" SDC sweep

$$y_m^{k+1} = y_{m-1}^{k+1} + \Delta t f(y_m^{k+1}) - \Delta t f(y_m^k) + \sum_{j=1}^M s_{m,j} f(y_j^k), \quad y_m^k \in \mathbb{R}^N$$


- Classically SDC is derived from error/correction equations (Dutt et al. 2000)

## Multi-level spectral deferred corrections (MLSDC)



Speck, Ruprecht, Emmett, Minion, Bolten, Krause

BIT Numerical Mathematics 2014



**Figure:** Hierarchy of three levels with 5 (red), 3 (blue) and 2 (black) collocation nodes and a coarsened spatial mesh for MLSDC applied to a semi-discrete time-dependent PDE.

- SDC sweeps on a **hierarchy** of levels  
→ SDC sweep analogue to relaxation in space MG
- Employ **FAS correction** to accurately represent solution on  $l > 1$
- Reduce overhead from  $f$  evaluations on higher levels by **coarsening**:  
→ reduced number of points, reduced order, reduced implicit solves, reduced physics, ...




# PFASST



Emmett, Minion

CAMCOS 2012



**Figure:** PFASST performs MLSDC cycles concurrently on multiple time steps, sending updated initial values after every sweep.

## PFASST in massively parallel simulations

- ▶ Particle based Navier-Stokes solver on 262,144 cores<sup>1</sup>
- ▶ Mesh-based combination with parallel multi-grid on 458,752 cores<sup>2</sup>

<sup>1</sup>Speck et al., Supercomputing 2012

<sup>2</sup>Ruprecht et al., Supercomputing 2013

# Inexact SDC



Speck, Ruprecht, Emmett, Minion, Krause  
 LNCSE 2015

## Algorithm 1: Standard SDC sweep

```

while (sdc_residual > time_tolerance) do
  for (m = 1 ... #substeps) do
    while (pmg_residual > space_tolerance) do
      | Perform PMG V-cycle
    end
    Perform backward Euler step:  $y_{m-1}^k \rightarrow y_m^k$ 
  end
end
end
  
```

## Algorithm 2: Sweep with inexact SDC

```

while (sdc_residual > time_tolerance) do
  for (m = 1 ... #substeps) do
    // Incomplete implicit solve
    for (k = 1 ... NumV) do
      | Perform PMG V-cycle
    end
    Perform approximate backward Euler step:  $y_{m-1}^k \rightarrow y_m^k$ 
  end
end
end
  
```

# IPFASST convergence



Minion et al.

SISC 2015




Figure: Convergence of IPFASST for simple  $u_t = u_{xx}$  problem depending on number of V-cycles.

- Stage  $y_m^k$  becomes increasingly accurate initial guess when solving for  $y_m^{k+1}$ :


$$y_m^{k+1} = y_{m-1}^{k+1} + \Delta t f(y_m^{k+1}) - \Delta t f(y_m^k) + \sum_{j=1}^M s_{m,j} f(y_j^k), \quad y_m^k \in \mathbb{R}^N$$

# IPFASST strong scaling



Minion et al.

SISC 2015



**Figure:** Strong scaling of IPFASST with parallel multi-grid solver (PMG) on the IBM Blue Gene/Q JUQUEEN at JSC. All shown runs provide the same accuracy.

- ▶ Space parallelization through parallel multigrid (code by M. Bolten)
- ▶ IPFASST can provide speedup beyond saturation of spatial parallelization alone

# Summary

- SDC is a preconditioned fixed point iteration to solve collocation equation
- MLSDC is multi-level extension with sweeps in a sense analogue to relaxation
- PFASST parallelizes MLSDC in time by iterating concurrently on multiple time-steps
- Iterative structure of SDC provides increasingly accurate starting values for PMG
  - can use approximate solves in implicit sub steps
  - corresponds to different order of time and space sweeps/cycles
- Results in ISDC, IMLSDC, IPFASST
- IPFASST shows good parallel performance on  $\mathcal{O}(10k - 100k)$  cores in first benchmarks

The End – for now.